

Autoencoders and Denoising Autoencoders

Lovedeep Gondara

MLRG, March 2018

Multilayer NN




Figure: Generic Neural Network

Multilayer NN

- ▶ We have labelled training examples
- ▶ Train the network to minimize output error (predicted vs real)
- ▶ i.e. Supervised training

Autoencoder

- ▶ We don't have training labels, just a dataset, unsupervised?
- ▶ What if we use the input as the labels
- ▶ Network tries to reconstruct the input, self-supervised?

Autoencoder




Figure: A simple autoencoder

Autoencoder

- ▶ Trying to learn an approximation to the identity function
- ▶ Identity function: $f(x) = x$
- ▶ Seems trivial, why do we care?

Autoencoder

- ▶ We can learn some interesting things from intermediate representations
- ▶ What if the number units in hidden layer are smaller than input?
- ▶ What if the number of units in hidden layer are larger than input?

Autoencoder

- ▶ We can learn some interesting things from intermediate representations
- ▶ What if the number units in hidden layer are smaller than input?
- ▶ What if the number of units in hidden layer are larger than input?

Autoencoder

- ▶ Can stack individual autoencoders to create a stacked deep autoencoder
- ▶ Yield better insights for large data
- ▶ Layerwise pre-training

Autoencoder

- ▶ Application to NLP: Word embeddings, Machine translation, Document clustering etc.
- ▶ Generic: Dimensionality reduction, Compression, Anomaly detection, pretraining etc.
- ▶ Image: Inpainting

Denoising Autoencoder

- ▶ Same as a generic autoencoder
- ▶ Input is corrupted
- ▶ Reconstruction is based on non-corrupted data

Denoising Autoencoder




Figure: Denoising example

Denoising Autoencoder

- ▶ Force the hidden layer to learn more robust features about the input while dealing with the noise/corruption
- ▶ i.e. not just an identity function
- ▶ Denoising autoencoder is a stochastic version of autoencoder

Denoising Autoencoder

- ▶ Tries to simultaneously encode the input and remove noise
- ▶ Noise is added stochastically at the input layer
- ▶ Different angles to understand the working of denoising autoencoder (manifold perspective, information theoretic, generative etc.)

Denoising Autoencoder




Figure: Denoising autoencoder

Denoising Autoencoder

- ▶ Robust representation: Learns from multiple input data representations
- ▶ To predict any subset of variables from the rest is a sufficient condition for completely capturing the joint distribution between a set of variables

Denoising Autoencoder

- ▶ Application to imaging: Denoising/completion
- ▶ Also can do most of the stuff an autoencoder can

Autoencoder/denoising autoencoder

- ▶ Generic design
- ▶ Can be constructed with any type of NN (CNN, RNN etc.)

References I



Le, Quoc V

A tutorial on deep learning part 2: autoencoders, convolutional neural networks and recurrent neural networks.

Google Brain, (2015): 1-20.



Goodfellow, Bengio and Courville

Deep learning book, Chapter 14: Autoencoders

MIT press, (2016)



Stanford deep learning course

<http://ufldl.stanford.edu/tutorial/unsupervised/Autoencoders/>